

To prove by induction that `PathFinder` computes the shortest paths and shortest-path distances from the source vertex `s` to each vertex `v` in the graph, we need to show that the breadth-first search (BFS) implemented in `PathFinder` correctly computes the shortest path for every vertex reachable from `s`.

Base Case (Step 0 of the BFS):

For the base case, consider vertices at distance 0 from the source.

- The only vertex at distance 0 is the source vertex itself.
- The source vertex `s` is enqueued into the `queue`, and its distance is set to `0`.
- In the constructor, we initialize `dist.put(s, 0)` for the source vertex `s`.
- This correctly represents that the distance to the source is 0, and the path to itself is trivial.
- **BFS Initialization:**
 - `queue.enqueue(s);`
 - `dist.put(s, 0);`

At this point, the only vertex in the queue is `s`, and the shortest distance from `s` to itself is correctly computed as `0`.

Thus, the base case holds: the shortest path from `s` to itself is correctly initialized to 0, and no other vertices are in the queue.

Inductive Hypothesis:

Assume that after $k \geq 0$ steps of BFS, the algorithm has correctly computed the shortest paths and shortest-path distances from the source `s` to all vertices that are at distance `k` from `s`.

Inductive Step (Step $k+1$ of the BFS):

We need to show that after $k + 1$ steps of the BFS, the algorithm correctly computes the shortest paths and distances for all vertices at distance $k + 1$ from `s`.

1. At the beginning of the $(k+1)$ -th step:

- All vertices at distance `k` from `s` have been dequeued from the `queue`, and for each of these vertices, all its adjacent vertices that have not yet been visited (i.e., not in `dist`) have been enqueued. For these enqueued vertices, their distance from `s` has been set to $k + 1$.

- This is guaranteed by the following code inside the BFS loop:

```

for (String w : G.adj(v)) {
    if (!dist.contains(w)) {
        queue.enqueue(w);
        dist.put(w, 1 + dist.get(v)); // Distance to w is set
to 1 + distance to v
        prev.put(w, v);                // Track the previous
vertex
    }
}

```

Consider a vertex v that is exactly $k+1$ edges away from the source s . Let u be the vertex that precedes v on the shortest path from s to v . By our inductive hypothesis, u is k edges away from s , and its shortest path and distance have been correctly computed.

2. Distance Computation:

- When the BFS reaches u , it explores all of u 's neighbors, including v .
- If v hasn't been visited yet (`!dist.contains(w)`), it sets:
`dist.put(w, 1 + dist.get(v));`
- This computes the distance to v as 1 plus the distance to u .
- Since u is on the shortest path to v , and the graph is unweighted, this distance is correct.

3. Path Computation:

- When v is discovered through u , the algorithm sets:
`prev.put(w, v);`
- This correctly records u as the predecessor of v on the shortest path.
- The `pathTo` method then reconstructs this path by following the `prev` links.

4. Optimality:

- BFS explores vertices in order of their distance from the source.
- When v is first discovered, it must be through a shortest path, as any longer path would have been explored later.
- Once v is added to the queue, its distance is set and never changed, ensuring the first (shortest) path is preserved.

Therefore, for vertices at distance $k+1$, the algorithm correctly computes both the shortest path and the shortest-path distance.

2. BFS Property:

- BFS explores all vertices at the same distance before moving on to vertices at the next greater distance. This ensures that when a vertex w is enqueued, it is the first time it is encountered, and its distance from s is the shortest path length. Once a vertex is dequeued, it is fully processed, meaning the shortest path to that vertex has already been found and no shorter path will be discovered.

- The algorithm does not re-enqueue any vertex once its distance has been set, which ensures that the first time a vertex is processed, the shortest path from `s` to that vertex is found.

Thus, after `k + 1` steps, all vertices at distance `k + 1` from `s` will be correctly processed, and their shortest distances will be set.

Conclusion (Inductive Proof):

By induction, after the BFS completes, the algorithm has computed the shortest path and shortest-path distance from the source vertex `s` to every vertex reachable from `s`. This holds because:

- The BFS processes vertices level by level (i.e., by increasing distance from the source).
- The first time a vertex is encountered in the BFS, the path leading to it is the shortest possible path.

Thus, the `PathFinder` class correctly computes the shortest paths and distances from the source `s` to each vertex in the graph.